

# ECS Fargate Spot Auto-Recovery - POC Design Document

## 1. Purpose of this POC

This Proof of Concept (POC) demonstrates an **automated, centralized recovery mechanism for ECS Fargate Spot workloads across multiple AWS accounts**.

The goal is to ensure **service continuity** when Fargate Spot capacity becomes unavailable, **without manual intervention** and **without writing per-account custom code**.

---

## 2. Problem Statement

ECS services running on **Fargate Spot** may experience:

- Capacity unavailability
- Spot interruptions
- Implicit scaling behavior without explicit failure signals

AWS does **not always emit reliable failure events** for these scenarios. As a result:

- ECS may silently rebalance tasks
- Spot failures may not trigger expected alarms
- Manual intervention is often required to switch to On-Demand

This POC solves the problem by:

- Centralizing detection and response
- Automatically failing over to On-Demand capacity
- Restoring Spot usage when capacity becomes stable

---

## 3. High-Level Architecture

### Key Design Principles

- **Centralized automation**
- **Zero per-account code changes**
- **Event-driven architecture**
- **Multi-account scalability**
- **Idempotent and safe operations**

### Account Structure

#### Central Account

Hosts:

- EventBridge (central bus)
- SQS queues
- Lambda automation
- DynamoDB state table

## Workload Account(s)

Host: - ECS clusters and services - Fargate + Fargate Spot capacity providers - Minimal onboarding configuration only

---

## 4. Workload Account Onboarding Template

Each new workload account must perform **only the following three steps:**

### Step 1: Create IAM Role

Create a role named:

CentralSpotAutomationRole

Trust policy allows the **central account** to assume this role.

### Step 2: Attach Permissions

Minimum permissions: - ecs:DescribeServices - ecs:UpdateService - iam:PassRole (if required)

### Step 3: Forward ECS Events

Create an EventBridge rule that forwards ECS events to the **central event bus ARN**.

No Lambda deployment  No DynamoDB  No per-service configuration

---

## 5. Central Account Components

**Important:** The entire Central Account setup for this POC is deployed using **AWS CDK**. No manual resource creation is required in the Central Account. The CDK stacks provision EventBridge rules, SQS queues, Lambda functions, DynamoDB tables, IAM roles, and permissions in a repeatable and auditable manner.

### 5.1 EventBridge

- Central event bus receives ECS events from all workload accounts
- Supports both real and synthetic test events

## 5.2 SQS Queues

Queue	Purpose
fargate-spot-placement-failure-queue	Placement failure events
fargate-spot-interruption-queue	Spot interruption events

Each queue has a Dead-Letter Queue (DLQ).

---

## 6. Lambda Functions

All Lambda functions in the Central Account are deployed and managed using **AWS CDK**. Each Lambda has explicitly scoped IAM permissions attached via CDK constructs, ensuring least-privilege access and consistent configuration across environments.

### 6.1 Placement Failure Handler (Python)

Triggered by: SQS

Responsibilities: - Parse ECS event - Extract account ID, region, cluster, service - Assume workload account role - Capture current Capacity Provider strategy - Store state in DynamoDB - Switch service to **On-Demand only**

Fallback Strategy Applied: - FARGATE\_SPOT → weight 0 - FARGATE → weight 10

---

### 6.2 Spot Interruption Handler (Node.js)

Triggered by: SQS

Responsibilities: - Handle real Spot interruption warnings - Apply failover logic - Ensure service continuity

---

### 6.3 Restore Handler (Python)

Triggered by: Scheduled EventBridge rule

Responsibilities: - Scan DynamoDB for pending restores - Assume workload account role - Restore original capacity provider strategy - Mark restore as complete - Emit audit event

---

## 7. DynamoDB State Table

### Table Name

```
ecs_cp_restore_state
```

### Primary Key

```
id = {account_id}::{cluster_name}::{service_name}
```

### Attributes

- config → Original capacity provider strategy
- event → Original triggering ECS event
- scheduled → true / false

This design guarantees: - Multi-account isolation - Idempotent updates - Safe retries

---

## 8. Event Handling & Testing Strategy

### Why Synthetic Testing Is Required

In real ECS behavior, Spot capacity loss does **not always emit explicit placement failure or interruption events**. ECS may simply reschedule or rebalance tasks automatically.

Because of this, for the POC we rely on **synthetic ECS events** to validate the automation end-to-end.

---

### Synthetic Event Injection (POC Testing)

Use the following command from **any AWS account that has permission to put events to the central event bus**:

```
aws events put-events
--region ap-south-1
--entries '[
{
  "Source": "test.aws.ecs",
  "DetailType": "ECS Service Action",
  "Resources": ["arn:aws:ecs:ap-south-1:903558039761:service/test-cluster2/
test-cluster-service"],
  "Detail": "{\"eventName\":\"SERVICE_TASK_PLACEMENT_FAILURE\", \"reason\":"
}
```

```
\\"RESOURCE:FARGATE\\",\\"capacityProviderArns\\":[\"arn:aws:ecs:ap-south-1:903558039761:capacity-provider/FARGATE_SPOT\"]}]'
```

This event simulates a Fargate Spot placement failure and triggers the complete pipeline:

EventBridge → SQS → Lambda → DynamoDB → ECS Update

---

## Verification Commands

After sending the synthetic event, validate each stage:

### 1. Check SQS queue depth

```
aws sqs get-queue-attributes  
--region ap-south-1  
--queue-url <QUEUE_URL>  
--attribute-names ApproximateNumberOfMessages  
ApproximateNumberOfMessagesNotVisible
```

### 2. Check Lambda execution logs

```
aws logs tail /aws/lambda/fargate-spot-task-placement-failure-handler  
--region ap-south-1  
--since 5m  
--follow
```

### 3. Verify DynamoDB state entry

```
aws dynamodb scan  
--region ap-south-1  
--table-name ecs_cp_restore_state
```

---

Successful execution confirms that the Central Account automation is functioning correctly without any dependency on real Spot failures.

## Synthetic Testing (POC)

Because ECS does not always emit failure events, synthetic events are used:

- `aws events put-events`
- Test placement failures
- Validate EventBridge → SQS → Lambda → DynamoDB flow

## Verification Steps

- Check SQS queue depth
  - Check Lambda CloudWatch logs
  - Scan DynamoDB for stored state
- 

## 9. Restore Scheduling

Restore Lambda runs on a schedule:

Examples: - Every 1 minute (DEV) - Every 5 minutes (PROD)

This ensures Spot capacity is reused when stable.

---

## 10. Security & IAM Model

- Central account uses **STS AssumeRole**
  - No long-lived credentials
  - No inbound access to workload accounts
  - Least-privilege policies enforced
- 

## 11. Scalability & Extensibility

This design supports: - Unlimited workload accounts - Multiple regions - Future extensions (notifications, metrics, dashboards)

No changes required in central automation when adding new accounts.

---

## 12. POC Success Criteria

The POC is considered successful if: - Spot failure triggers automated failover - Services continue running on On-Demand - Original strategy is restored automatically - No manual intervention is required - New workload accounts onboard with only 3 steps

---

## 13. Conclusion

This POC demonstrates a **robust, scalable, and production-ready pattern** for managing ECS Fargate Spot reliability across multiple AWS accounts using centralized automation.

It eliminates operational toil, improves service availability, and provides a clean path toward enterprise-grade Spot adoption.